

Final Report

Distributed Ticketing System

DistribuTix

CPSC 559

Introduction to Distributed Systems

Winter 2024

Group 19

Redwanul Islam

Ehab Islam

Mahtab Khan

Zeyad Omran

Ahnaf Eusa

April 7, 2024



Contents

Contents	2
Preface	3
Abstract	4
Introduction	5
0.1 Background	5
0.2 Motivation and Objective	5
Architecture	6
1.1 Architecture (Overall Design)	6
1.2 Intended User Base	10
1.3 System Requirements	10
1.4 System Inputs	11
1.5 Workflow of Data	12
1.6 System Guarantees	13
Communication	14
2.1 RESTful APIs: Enabling Flexible Service Interaction	14
2.2 JSON: Streamlining Data Serialization	15
2.3 Layered Protocols: Foundation of Network Communication	15
2.4 Virtual Private Cloud (VPC): Enhancing Security and Control	16
Synchronization	17
3.1 The Ring Algorithm	17
3.2 Multithreading with Locking Mechanism	18
3.3 Concurrency and Consistency	18
Replication	19
4.1 Implementing Passive Replication Strategy	19
4.2 Redundancy and Failover	19
Consistency	21
5.1 Ensuring Consistency in Our Authentication and Order Processing System	21
5.2 Multithreading and Locking Mechanism	21
5.3 Achieving Eventual Consistency	21
Fault Tolerance	23
6.1 Service Replication and Failover Mechanisms	23
6.2 Comprehensive Failure Detection and Management	24
Limitations & Challenges	25
7.1 Challenges with Ring Algorithm	25
7.2 Limited Dynamic Scalability	25
Summary & Conclusions	26

Preface

This report is an amalgamation of all the high-yielding results of our hard work over the last four months.

For its inception, we would like to thank Professor Kawash and our diligent TAs for structuring the course, outlining the project, and putting in all the effort towards the objective.

Most of us in this group are in our last semester. And we have known each other mostly since the first year. Thank you for giving us a reason to be able to work together, learn, and create something together in the process.

DistribuTix is the answer we have produced for the project outline; showcasing all that we have learned over the course and supplementary individual experiences we have gained over the years.

We had a lot of fun working on this project.
We hope you have some too while reviewing this.

- Group 19

Abstract

The final report on **DistribuTix**, presented by Group 19, unveils a distributed ticketing system designed to transform the traditional ticketing landscape through advanced distributed systems principles. This system prioritizes scalability, reliability, and an enhanced user experience, employing RESTful APIs and JSON for efficient data serialization, alongside sophisticated replication strategies to ensure system resilience and data integrity. Key to its design is the implementation of leader election algorithms and passive replication mechanisms, which underpin the system's fault tolerance and consistency. The architecture further integrates a robust locking mechanism within its multithreading environment, safeguarding against concurrent data access issues and maintaining transactional integrity across the system.

DistribuTix stands as a testament to the potential of distributed systems in addressing the complexities of real-time ticketing demands. The system's innovative approach to leader election, combined with its strategic replication and consistency models, ensures a reliable and seamless user experience, even under peak loads. By incorporating these technologies, DistribuTix not only showcases a leap forward in ticketing solutions but also provides a blueprint for future advancements in distributed system applications. This project highlights the team's successful endeavor in crafting a system that is both scalable and robust, ready to meet the dynamic needs of the ticketing industry.

Introduction

0.1 Background

The development of distributed systems has revolutionized the way software applications operate, enhancing scalability, reliability, and performance across various industries. Within this context, our project focused on creating a distributed ticket system, aimed at addressing the complexities and inefficiencies inherent in traditional ticketing systems. Such systems are pivotal for numerous sectors, including transportation, entertainment, and customer service, where the need to manage large volumes of ticket transactions simultaneously is critical.

0.2 Motivation and Objective

Our motivation for developing the distributed ticket system stems from a desire to resolve the inefficiencies and limitations of traditional centralized ticketing frameworks, particularly in handling high-demand scenarios. The key objectives of our project were carefully chosen to address these challenges and revolve around five pivotal areas: consistency, user experience, reliability, availability, and scalability.

We aimed to ensure real-time data consistency across the system, particularly for ticket availability, to avoid overbooking and discrepancies. A smooth and rapid ticket purchasing process was prioritized to enhance the user experience, reducing wait times and simplifying transactions. The system's design focuses on reliability to maintain stability and consistent performance, even under varying operational conditions. By employing a fault-tolerant architecture, we sought to guarantee the system's availability around the clock, ensuring users could conduct transactions at any time. Finally, the system was engineered for scalability to efficiently manage an increase in user requests during peak times, ensuring it could accommodate large volumes of transactions without degradation in performance. These objectives collectively underpin our system's design, aiming to offer a robust, user-friendly ticketing solution that can adapt to future demands.

Architecture

1.1 Architecture (Overall Design)

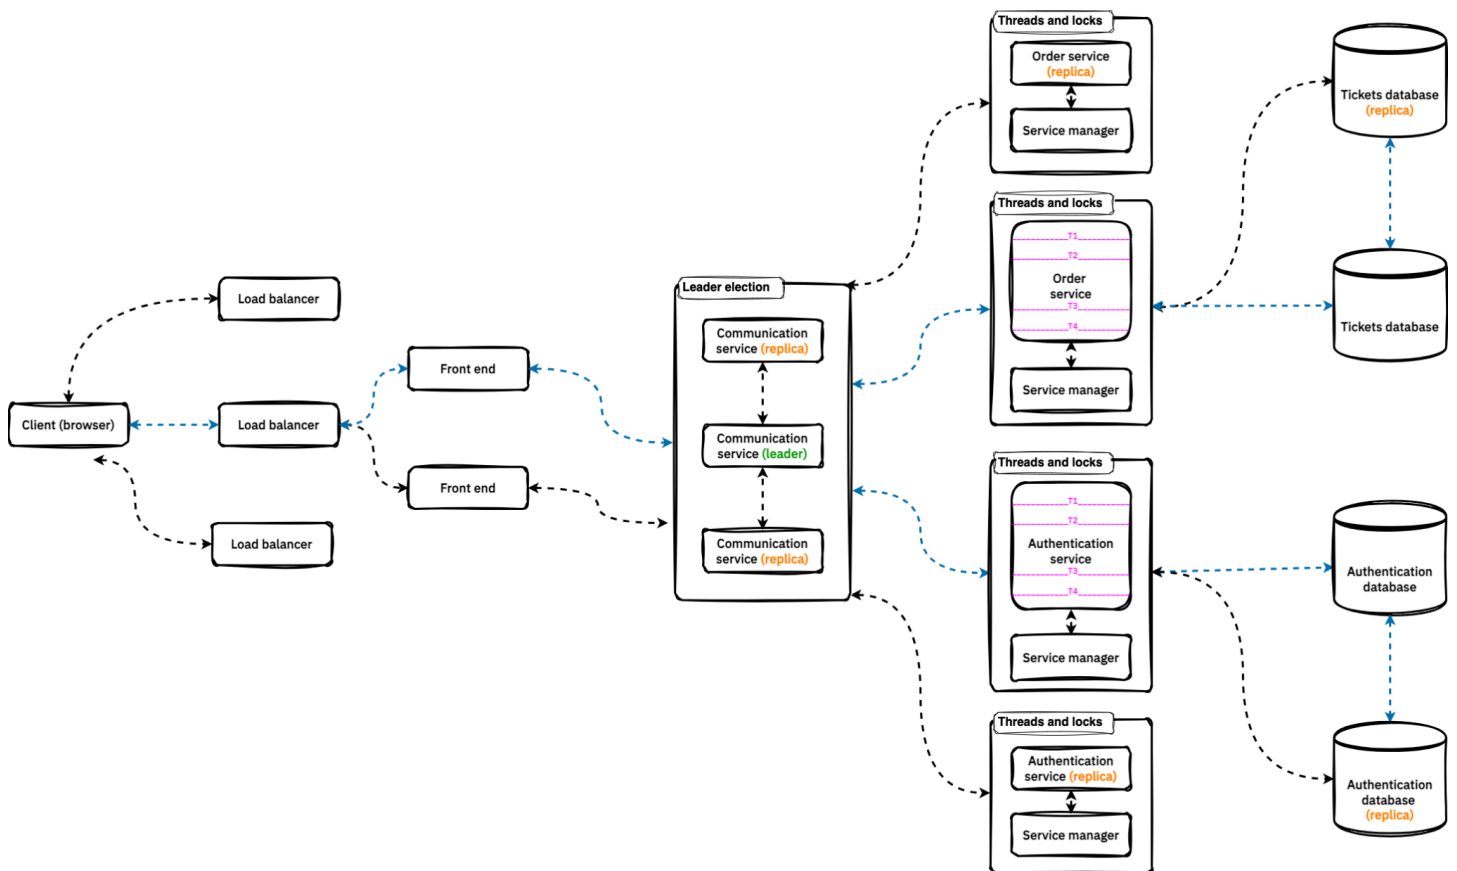


Figure 1. High-level architecture

The architecture diagram illustrates a distributed system designed for high availability and scalability. Here's an explanation of its components and their interactions in the following pages.

Client (Browser)

This represents the user's point of interaction with the system via a web browser. There can be multiple clients trying to interact with the load balancer.

Load Balancer

Serves as the traffic controller that distributes incoming client requests evenly across the frontend servers to balance the load and ensure no single frontend server becomes a bottleneck. The load balancer is replicated in 3 availability zones by AWS.

Front End

The servers in question facilitate direct interaction with the client's browser via the load balancer, delivering the user interface and processing incoming requests. To ensure uninterrupted service and resilience, an additional front-end server is operational as a standby. This server is poised to seamlessly take over request handling should the primary service become unavailable, thus maintaining system redundancy and reliability.

Communication Service

The communication service handles internal communication between different services in the architecture. It has a leader process and replicas for fault tolerance.

The leader election process ensures that there is always one leader to coordinate actions among replicas, thus providing consistency and coordination across services.




Figure 2. Communication service

Order Service and it's Replica

The order service functions as a dedicated microservice, orchestrating the processing of orders. It is complemented by an active standby replica, enhancing the system's scalability by facilitating independent adjustments to service capacity as needed.

As a multithreaded service, it employs synchronization mechanisms, including locks, to ensure data consistency and integrity during concurrent operations.




Figure 3. Order service and its standby replica

Authentication Service and it's Replica

The authentication service functions as a dedicated microservice, handling user authentication. It is complemented by an active standby replica, enhancing the system's scalability by facilitating independent adjustments to service capacity as needed.

As a multithreaded service, it employs synchronization mechanisms, including locks, to ensure data consistency and integrity during concurrent operations.




Figure 4. Authentication service and its standby replica

Service Manager

This is a central control system or service registry that keeps track of all the services, their instances, and their health status.

Databases

There are separate databases for tickets and authentication, each with its replica. The use of replicas indicates a need for high availability and data redundancy.

InterService Communication

The dashed blue lines (*in Figure 1*) represent the flow of requests and data for the active services/databases.

It represents the communication between the different services and their respective databases, showing that these services work independently but may need to communicate or coordinate for some operations.

The dashed black lines represent the flow of potential requests and data for the standby services/databases.

It has the same meaning as the dashed blue lines except it is for standby servers.

1.2 Intended User Base

Primary Users: The system's versatile architecture allows for a multitude of users.

Event Organizers and Venues: Coordinators of live events such as concerts, sports competitions, theatrical productions, and conferences, this system is indispensable for the real-time issuance and validation of tickets, ensuring a seamless entry experience for attendees.

Event Attendees: Individuals who visit these events and venues benefit from a reliable and efficient ticketing process. The system's capability to handle high demand and provide quick responses enhances the overall customer experience, from purchase to admission.

Additional Users: The system's versatile architecture allows for customization to suit a variety of other applications and user groups. These may include:

Cinema Chains: Designed to meet the demands of multiplexes and cinema chains, the system facilitates the smooth booking of movie tickets, accommodating advanced reservations, and on-the-spot purchases, enriching the cinema-goer's experience.

1.3 System Requirements

- At least one host for each type of service: Communication, Authentication, Order service, and Front-end (at least 2 recommended for failover and high availability)
- Private network between the physical nodes
- SQL database (RDS recommended)
- Windows

1.4 System Inputs

The web application provides a user-friendly interface that guides users through straightforward input fields. Below is a detailed breakdown of the inputs required on various pages within the system:

Authentication Page

Users will enter their **Username** and **Password** to either sign up for a new account or log in to an existing account.

Create Ticket Page

Here, users can list a new event by entering the **Event Name**, **Description**, **Price**, **Date of the Event**, **Location Name**, **Coordinates**, and **Quantity** of tickets available.

Two buttons are present: one to **Cancel** the creation process and another to **Create Ticket** confirming the event details.

Home Page

Users can filter events by specifying a **Start Date** and **End Date**.

A search field allows for entering a **Search** description to find specific events.

A prominent **Create Ticket** button is accessible on the top right-hand side of the page for quick navigation to event creation.

Ticket Detail Page

Two interactive buttons are available: one to **Order Now**, initiating the purchase process, and another to **Go Back**, returning to the previous view.

This structure ensures that users can interact with the application efficiently, whether they are creating, searching for, or purchasing tickets.

1.5 Workflow of Data

Client to Load Balancer to Front End

When a client (user's browser) makes a request, it first hits the load balancer. The load balancer then directs the request to the active front-end server.

Front End to Communication Service

The front-end server processes the request and, if necessary, forwards it to the communication service. For instance, if the request is related to user authentication, it's sent to the Authentication Service. If it's related to ticket management, it goes to the Order Service.

Communication Service to Authentication/Order Service

Case 1:

If the client's request is related to user authentication, it's sent to the active Authentication Service.

Case 2:

If the request is related to ticket management, it goes to the active Order Service.

Service to Database

If necessary, the Order Service interacts with the Tickets Database, while the Authentication Service interacts with the Authentication Database. Data is read from or written to these databases as part of the service's operations.

Feedback Loop

- Once the backend services process the user's request, the outcome is communicated back to the front end through the same pathway. The front end, in turn, conveys this response to the client's browser.
- Upon completing registration or login, users are redirected appropriately, affirming the successful execution of their requested operations. Additionally, users are presented with a confirmation of their actions, such as successful ticket purchase or reservation details. This feedback is integral for a seamless user experience, keeping the user informed at each step of the interaction with the system.

The overall data flow supports a robust, fault-tolerant, and scalable system, capable of handling a significant number of concurrent operations while maintaining data integrity and availability.

1.6 System Guarantees

Scalability

The system is designed to efficiently handle an increasing number of user requests, especially during peak times when demand surges, such as immediately after ticket sales open for a popular event. Scalability ensures that the system can expand its capacity to manage higher loads without performance degradation, which is often achieved through the addition of resources or the distribution of load across multiple servers.

Availability

High availability is crucial for the ticketing system to ensure that it is always online and accessible for users to conduct transactions. This means implementing failover strategies, redundancy, and robust error-handling mechanisms so that system downtime is minimized and users can reliably access the service whenever they need to.

Reliability

The system must maintain stability and deliver consistent performance over time, regardless of the number of users or the complexity of transactions. This means that the system must be resilient to failures, bugs, and errors, recovering gracefully from any issues and providing a reliable service that users can trust.

User Experience

A seamless user interface and a smooth, fast ticket-purchasing process are essential for customer satisfaction. The system must be intuitive and straightforward, minimizing the time and effort required for users to find and purchase tickets. This involves optimizing workflows, reducing latency, and ensuring that the user experience is pleasant and efficient from start to finish.

Consistency

Ensuring real-time data uniformity is critical, especially in scenarios like ticket availability where multiple users may be attempting to purchase a limited number of tickets simultaneously. The system must reflect the most current state of data across all nodes so that every user sees the same information regarding ticket availability, prices, and event details, preventing overbooking and data conflicts.

Communication

Our system incorporates a blend of sophisticated technologies and architectural principles to facilitate secure, reliable, and efficient communication across its diverse components. By leveraging RESTful APIs, Layered Protocols, a Virtual Private Cloud (VPC), and JSON for data serialization, **DistribuTix** ensures not only fault tolerance and scalability but also a seamless and responsive user experience.

Here, we explore how each of these technologies—RESTful APIs, JSON, Layered Protocols, and VPC—plays a crucial role in **DistribuTix**'s architecture, each contributing uniquely to the system's overall functionality.

2.1 RESTful APIs: Enabling Flexible Service Interaction

RESTful APIs stand at the core of **DistribuTix**'s service interaction model, facilitating stateless communication between the system's frontend and backend services, as well as among the backend services themselves.

- **Design Principles:** These APIs are designed around the principles of REST, using standard HTTP methods (GET and POST) to operate on resource representations. This stateless approach simplifies the architecture, making it more scalable and tolerant of faults.
- **Implementation:** **DistribuTix** utilizes Flask, a micro web framework, to implement these RESTful services. Flask's minimalistic yet powerful toolkit enables rapid development and deployment of API endpoints, catering to both simple and complex service logic requirements.

2.2 JSON: Streamlining Data Serialization

JSON (JavaScript Object Notation) is the lifeblood of data interchange within **DistribuTix**, enabling the efficient exchange of data across services.

- **Application:** Every RESTful API call in **DistribuTix** leverages JSON to encapsulate request data and responses. This consistency in data format simplifies the processing and validation of client-server interactions, enhancing the system's overall responsiveness and reliability.
- **Advantages:** Its lightweight nature, readability, and ease of use in web technologies make JSON an ideal format for web-based applications like **DistribuTix**. It seamlessly integrates with our JavaScript-based frontend and is easily parsed by backend languages such as Python, which is used extensively in **DistribuTix**.

2.3 Layered Protocols: Foundation of Network Communication

The use of Layered Protocols in **DistribuTix** ensures efficient and secure data transmission across the system's distributed components. This hierarchical organization of network protocols into layers abstracts the complexities of network communication, allowing for modular and flexible system design.

Protocol Stack: At the heart of **DistribuTix**'s networking is the TCP/IP suite, with HTTP/HTTPS protocols facilitating application-layer communications, TCP ensuring reliable transport, and IP handling routing. This stack is crucial for the integrity and security of data exchange within **DistribuTix**.

Benefits: The abstraction provided by layered protocols enables **DistribuTix** to easily adapt to new technologies and scale with minimal adjustments to the overarching system. It ensures interoperability across different network technologies and streamlines the integration of new services.

2.4 Virtual Private Cloud (VPC): Enhancing Security and Control

DistribuTix's deployment within an AWS Virtual Private Cloud (VPC) provides a secure, isolated network for all system communications, crucial for protecting sensitive data and operations.

Network Isolation: The VPC creates a segregated environment for **DistribuTix**, isolating its infrastructure from the broader internet. This isolation is pivotal in minimizing the risk of unauthorized access and exposure to potential security threats, ensuring that sensitive data and critical operations are shielded within a controlled network space.

Lower Latency and Faster Response Times: Operating within a VPC also presents significant performance benefits. By optimizing the network paths and reducing the distance that data must travel between **DistribuTix** components and its users, the system achieves lower latency and faster response times. This is particularly beneficial for a ticketing system where timely processing and responsiveness are crucial for user satisfaction.

Added Security: With the integration of VPC security groups, **DistribuTix** gains an additional layer of security. This allows for granular control over inbound and outbound traffic, ensuring that only legitimate and authorized requests are processed. This meticulous approach to network traffic management further strengthens the system's defense against potential security breaches.

Synchronization

To maintain the integrity and consistency of data within our **DistribuTix** system, we have implemented a sophisticated synchronization mechanism underpinned by a ring algorithm and supported by a multithreading model with an accompanying locking mechanism.

3.1 The Ring Algorithm

Our choice of a ring algorithm is grounded in its ability to lower performance overheads, making it suitable for environments where a limited number of processes are in play. The simplicity of this algorithm lies in its less complex structure, as the leader is elected based on straightforward attributes like the port number, which mitigates the need for intricate coordination. This simplicity translates into increased speed for certain critical operations, particularly those associated with synchronization and consistency.

The ring structure facilitates the efficient passing of messages or tokens, allowing each node in the system to operate with a clear understanding of its predecessor and successor. This clarity ensures that data remains consistent across the system with minimal replication lag. It also simplifies the management of concurrency as the ring algorithm inherently sequences operations, reducing conflict and contention.

Pseudocode:

- The **initiate_election** endpoint starts an election process by spawning a new thread that calls the **send_election_message** function.
- In **send_election_message**, the current process sends an election message to the next process in the ring. If a request fails, it tries the next process in the sequence based on the port number.
- The **election** endpoint processes incoming election messages. If the message has circled back with the same service ID, this process declares itself the leader and announces it to the other processes. If the incoming ID is higher than the process's ID, it forwards the message. If it's lower, it replaces it with its own ID and forwards it.
- The **send_leader_announcement** function sends out an announcement to all processes that a new leader has been elected.

3.2 Multithreading with Locking Mechanism

Within our multi-threaded environment, consistency is further reinforced through the use of a locking mechanism. When our system is presented with write operations, these locks ensure that only one thread can execute a write SQL statement at a time, thereby preventing race conditions and maintaining data integrity. Worker thread assignment is carefully managed so that each new task, such as an incoming order, is handled efficiently without compromising the responsiveness of the system.

Our implementation of eventual consistency within this model accommodates the realities of a distributed system. It allows for a period wherein not all replicas may be perfectly synchronized but ensures that, over time, all states converge consistently. This approach not only guarantees data consistency in the long term but also allows the system to handle a high volume of concurrent operations without significant performance degradation.

3.3 Concurrency and Consistency

The interplay between concurrency and consistency is delicately balanced in our system. While concurrency aims to maximize system utilization and throughput, consistency ensures that all operations on data yield correct and predictable results. By carefully managing these two aspects, **DistribuTix** achieves a state where operations are performed concurrently without sacrificing the consistency of the data.

The diagrams included in the slides visually represent the ring algorithm's topology and the locking mechanism symbolically, providing a clear illustration of how these components function within our system to maintain synchronization and consistency.

In conclusion, our **DistribuTix** system's approach to synchronization and consistency is tailored to manage the complexities associated with distributed systems. By combining the efficiency of the ring algorithm with the robustness of a multithreaded locking mechanism and the reliability of eventual consistency, we ensure that our system remains both performant and consistent, ready to meet the demanding needs of our clients.

Replication

4.1 Implementing Passive Replication Strategy

Our system leverages a passive replication strategy to enhance consistency across its services. In this model, replicas remain in a ready state but do not actively handle requests under normal operation. They are on standby, prepared to take over immediately should the primary service become unavailable.

This passive approach to replication ensures a high level of consistency for several reasons. Firstly, by avoiding active handling of user requests by the replicas, the risk of state divergence between the primary and the replicas is minimized. The primary service handles all write operations, which when needed by the replica will already be committed to the Database. This synchronous replication means that at any point in time, the state of the primary and the standby replicas is consistent.

Secondly, passive replication serves as a foundational component for our robustness against data inconsistency. Should a primary service or database fail, the standby replica possesses an up-to-date and consistent state that is a direct mirror of the primary prior to its failure. Thus, passive replication supports our system's overall consistency by ensuring that replicas are always synchronized and ready to serve an accurate and consistent state without delay or the need for state reconciliation.

4.2 Redundancy and Failover

The system's architecture is designed with redundancy as a core principle, ensuring that for every critical component, there is a replica that maintains a consistent state. This redundancy is vital for our system's resilience, providing a safety net in case of service failure. At any given moment, the system is prepared with a consistent state that can be relied upon.

Failover mechanisms are integral to this redundant structure. They are designed to detect failures in real-time and initiate an automatic and seamless transition from the failed primary service to the standby replica. The key to this seamless switch is the guarantee that the standby is always in a consistent state with the primary. There is no interruption to the user experience during this transition, as the failover process ensures that the replica service is immediately available to handle requests. This immediate availability is critical to maintaining a consistent user experience and system reliability.

The passive replication and failover process are illustrated in our system diagrams, showcasing the flow from active services to their respective replicas. The diagrams indicate how each component fits into the overall system architecture, providing a visual understanding of how consistency is maintained through these mechanisms.

In conclusion, our system's design for passive replication and redundancy ensures a high degree of consistency and reliability. This design philosophy is fundamental in providing our users with a resilient, robust, and trustworthy experience, even in the face of potential system failures.

Consistency

5.1 Ensuring Consistency in Our Authentication and Order Processing System

In the development of our authentication and order processing system, we have meticulously integrated a multithreading model combined with an advanced locking mechanism. This design choice is pivotal in ensuring that our system processes only one write SQL statement at a time, thereby safeguarding the consistency and integrity of our data.

5.2 Multithreading and Locking Mechanism

Our system architecture is built to handle high volumes of transactions efficiently. Upon receiving a new order, the system assigns it to a dedicated worker thread. This allocation strategy is not only optimal for maximizing our processing efficiency but also plays a critical role in maintaining low response times for our users.

The cornerstone of our approach is the locking mechanism that accompanies the multithreading model. By enforcing exclusive access for write operations, we eliminate the risk of concurrent modifications leading to data anomalies or inconsistencies. This ensures that even in a highly threaded environment, transactions are executed in a controlled sequence, maintaining the integrity and consistency of our database at all times.

5.3 Achieving Eventual Consistency

Our adherence to the eventual consistency model is a deliberate choice, designed to offer both robustness and reliability in order processing. This model, while allowing for some degree of latency in achieving absolute consistency across all nodes, ensures that our system remains highly available and resilient to various operational stresses.

In practice, this means that our system is well-equipped to handle a substantial influx of requests without sacrificing the consistency and integrity of the data. The eventual consistency model provides us with the flexibility needed to optimize our system's performance and scalability, ensuring that all transactions eventually reach a consistent state across the system.

By leveraging a sophisticated blend of multithreading, locking mechanisms, and the eventual consistency model, our authentication and order processing system stands out as both efficient and reliable. This architecture allows us to process a high volume of transactions sequentially,

ensuring data consistency and maintaining the overall integrity of our system. Our commitment to these principles supports our goal of delivering unwavering consistency and integrity in order processing, making our system robust and ready to meet the demands of our users.

Fault Tolerance

Our system incorporates an advanced monitoring mechanism that conducts regular health checks on services. This proactive approach ensures any service going offline is immediately identified, allowing for swift restarts and minimizing operational disruptions. This process is facilitated by a service monitor that continuously evaluates the health and status of each component within our infrastructure.

Utilizing **REST API** calls as a heartbeat mechanism, the monitor verifies the operational status of services, ensuring they remain responsive and functional.

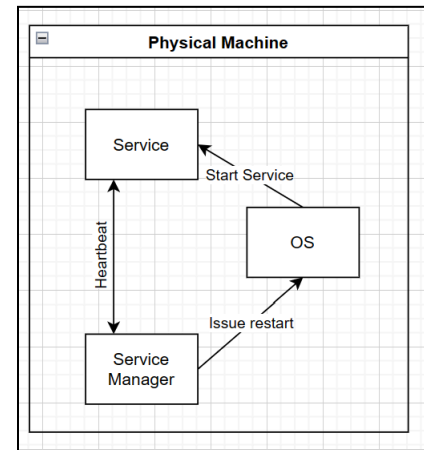


Figure 5. Heartbeat Mechanism

Our approach includes a distributed architecture, where each server hosting a service is equipped with a health monitoring service. This specialized service attempts to automatically restart any faulty service by terminating the unresponsive process and initiating a restart command. This automation enhances our system's capability to recover from failures promptly and reduces reliance on manual intervention, significantly cutting down recovery times and the potential for human error.

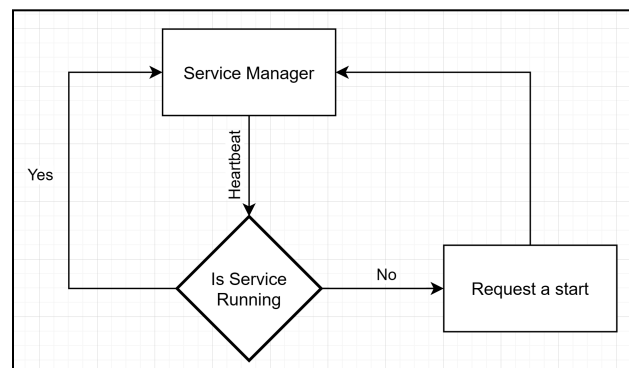


Figure 6. Heartbeat logic

6.1 Service Replication and Failover Mechanisms

To further bolster our system's resilience, we implement service replication. This strategy involves duplicating critical services, thereby allowing tasks to automatically switch to backup instances in the event of a primary service failure. This seamless transition ensures that users experience no interruption in service availability.

6.2 Comprehensive Failure Detection and Management

Our system is equipped to detect and manage a wide range of failure types, including:

Crash Failures

By running multiple instances of each service, we ensure that a single crash does not halt our operations. This redundancy is crucial for maintaining service continuity.

Omission Failures

In the event of failed requests, our system automatically reroutes these to operational services, ensuring uninterrupted service delivery.

Timing Failures

Our system is designed to handle delays efficiently. When delays are detected, requests are retried with replicated services to maintain speed and efficiency.

Response Failures

Transaction verification is a critical component of our failure management strategy. We ensure transactions are confirmed before completion, and in case of discrepancies, operations are reverted to maintain data integrity.

Arbitrary Failures

To mitigate security risks, our services communicate over a private network. This isolation helps reduce the likelihood of arbitrary failures impacting our system.

Limitations & Challenges

While the **DistribuTix** system has achieved significant milestones in addressing the complexities and inefficiencies of traditional ticketing systems through distributed technologies, it is not without its limitations and challenges. Two of the primary constraints encountered in the course of this project include the system's challenges with the ring algorithm in terms of scalability and its limited dynamic scalability.

7.1 Challenges with Ring Algorithm

The worst-case time complexity of $O(n^2)$ for the ring algorithm implies that as more processes are added, the time taken for leader election increases significantly. This scaling issue could lead to delays, adversely affecting the user experience. Therefore, considering alternatives like the Bully Algorithm, which might offer better performance characteristics in certain scenarios, becomes crucial. These alternatives could potentially streamline the leader election process, ensuring more efficient operation and a smoother experience as the system scales.

7.2 Limited Dynamic Scalability

Another significant challenge faced by **DistribuTix** is its limited capability for dynamic scalability. The current system architecture requires manual code adjustments to add more nodes to the network. This limitation hinders the system's ability to automatically scale its resources in response to real-time demand fluctuations. In high-demand scenarios, such as the sale of tickets for a highly anticipated event, the inability to dynamically scale could result in performance bottlenecks, increased latency, and a degraded user experience.

Summary & Conclusions

Our project embarked on addressing the inefficiencies and limitations inherent in traditional ticketing systems through the development of a distributed ticketing system named **DistribuTix**. This system was designed with the ambition of revolutionizing the ticketing landscape by ensuring scalability, reliability, and an enhanced user experience through the utilization of distributed systems principles.

Throughout the project, we successfully implemented an architecture that promotes high availability, fault tolerance, and consistency across all nodes in the system. By leveraging technologies such as RESTful APIs for seamless service interaction, JSON for efficient data serialization, and employing a mix of synchronous and asynchronous replication strategies, **DistribuTix** was engineered to withstand the rigors of peak demand scenarios while maintaining data integrity and operational stability. One of the key achievements of our project was the implementation of a robust fault tolerance mechanism that utilizes passive replication to ensure service continuity in the face of node failures. This, combined with our strategic approach to synchronization and consistency through a multithreading model and locking mechanisms, has set a new benchmark for reliability in distributed ticketing solutions.

However, our journey was not without its challenges. Ensuring real-time data consistency across all nodes in the face of network partitions and varying latencies posed a significant hurdle. Moreover, balancing the system's scalability with fault tolerance demanded meticulous planning and optimization to avoid performance bottlenecks. Looking ahead, there is substantial scope for enhancing **DistribuTix** by exploring more advanced distributed algorithms that could further optimize fault tolerance and data consistency. Additionally, integrating machine learning algorithms for predictive analytics could offer insights into user behavior, enabling anticipatory scaling and resource management. The incorporation of blockchain technology could also enhance security and transparency in ticket transactions, opening new avenues for trust and verification in the ticketing ecosystem.

In conclusion, **DistribuTix** represents a significant step forward in the realm of distributed ticketing systems, providing a scalable, reliable, and user-friendly solution. Our project's findings and methodologies contribute valuable insights into the design and implementation of distributed systems, offering a solid foundation for future research and development in this field.

In the development of this project, generative AI was used for preliminary research, idea generation, and text composition, providing insights into distributed systems and their applications in digital ticketing solutions.